

ResNet20

October 27, 2025

```
[1]: # HW

# 1. train resnet20 and vgg16 to achieve >90% accuracy
# 2. save your trained model in the result folder
# 3. Restart your jupyter notebook by "Kernel - Restart & Clear Output"
# 4. Load your saved model for vgg16 and validate to see the accuracy
# 5. such as the last part of "[W2S2_example2]_CNN_for_MNIST.ipynb", prehook ↵
#    the input layers of all the conv layers.
# 6. from the first prehooked input, compute to get the second prehooked input.
#    ↵
# 7. Compare your computed second input vs. the prehooked second input.
```

```
[2]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

from models import *    # bring everything in the folder models

global best_prec
use_gpu = torch.cuda.is_available()
print('=> Building model...')

batch_size = 128

model_name = "ResNet20"
model = resnet20_cifar()
```

```

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243, 0.262])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
    ]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
                                          shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
                                         shuffle=False, num_workers=2)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch includes "batch_size" data points
# CIFAR10 has 50,000 training data, and 10,000 validation data.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter() ## at the begining of each epoch, this should be reset
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time() # measure current time

```

```

for i, (input, target) in enumerate(trainloader):
    # measure data loading time
    data_time.update(time.time() - end) # data loading time

    input, target = input.cuda(), target.cuda()

    # compute output
    output = model(input)
    loss = criterion(output, target)

    # measure accuracy and record loss
    prec = accuracy(output, target)[0]
    losses.update(loss.item(), input.size(0))
    top1.update(prec.item(), input.size(0))

    # compute gradient and do SGD step
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # measure elapsed time
    batch_time.update(time.time() - end) # time spent to process one batch
    end = time.time()

    if i % print_freq == 0:
        print('Epoch: [{0}][{1}/{2}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  epoch, i, len(trainloader), batch_time=batch_time,
                  data_time=data_time, loss=losses, top1=top1))

```



```

def validate(val_loader, model, criterion):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

```

```

    input, target = input.cuda(), target.cuda()

    # compute output
    output = model(input)
    loss = criterion(output, target)

    # measure accuracy and record loss
    prec = accuracy(output, target)[0]
    losses.update(loss.item(), input.size(0))
    top1.update(prec.item(), input.size(0))

    # measure elapsed time
    batch_time.update(time.time() - end)
    end = time.time()

    if i % print_freq == 0: # This line shows how frequently print out
        ↪the status. e.g., i%5 => every 5 batch, prints out
        print('Test: [{0}/{1}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  i, len(val_loader), batch_time=batch_time, loss=losses,
                  top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg

def accuracy(output, target, topk=(1,5)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk) # 5
    batch_size = target.size(0) # 128

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None, ↪
    ↪largest=True, sorted=True) # will output (max value, its index)
    pred = pred.t() # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred)) # "-1": calculate ↪
    ↪automatically

    res = []
    for k in topk: # 1, 5
        correct_k = correct[:k].reshape(-1).float().sum(0) # reshape(-1): make ↪
        ↪a flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size)) # correct: size of ↪
        ↪[maxk, batch_size]

```

```

    return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n      ## n is impact factor
        self.count += n
        self.avg = self.sum / self.count

    def save_checkpoint(state, is_best, fdir):
        filepath = os.path.join(fdir, 'checkpoint.pth')
        torch.save(state, filepath)
        if is_best:
            shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

    def adjust_learning_rate(optimizer, epoch):
        """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120 epochs"""
        adjust_list = [150, 225]
        if epoch in adjust_list:
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['lr'] * 0.1

#model = nn.DataParallel(model).cuda()
#all_params = checkpoint['state_dict']
#model.load_state_dict(all_params, strict=False)
#criterion = nn.CrossEntropyLoss().cuda()
#validate(testloader, model, criterion)

=> Building model...
Files already downloaded and verified
Files already downloaded and verified

```

```
[3]: # import matplotlib.pyplot as plt
# import numpy as np

# # functions to show an image

# def imshow(img):
#     img = img / 2 + 0.5      # unnormalize
#     npimg = img.numpy()
#     plt.imshow(np.transpose(npimg, (1, 2, 0)))
#     plt.show()

# # get some random training images
# dataiter = iter(testloader)
# images, labels = next(dataiter) ## If you run this line, the next data batch is called subsequently.

# # show images
# imshow(torchvision.utils.make_grid(images))
```

```
[4]: # # This cell is from the website
# lr = 5e-3
# weight_decay = 1e-5
# epochs = 200
# best_prec = 0

# fdir = 'result/'+str(model_name)+'/model_best.pth.tar'
# checkpoint = torch.load(fdir)
# model.load_state_dict(checkpoint['state_dict'])
# model = model.cuda()
# criterion = nn.CrossEntropyLoss().cuda()
# optimizer = torch.optim.SGD(model.parameters(), lr=lr, momentum=0.93, weight_decay=weight_decay)
# # weight decay: for regularization to prevent overfitting

# if not os.path.exists('result'):
#     os.makedirs('result')

# fdir = 'result/'+str(model_name)

# if not os.path.exists(fdir):
#     os.makedirs(fdir)

# for epoch in range(0, epochs):
#     adjust_learning_rate(optimizer, epoch)

#     train(trainloader, model, criterion, optimizer, epoch)
```

```

#     # evaluate on test set
#     print("Validation starts")
#     prec = validate(testloader, model, criterion)

#     # remember best precision and save checkpoint
#     is_best = prec > best_prec
#     best_prec = max(prec,best_prec)
#     print('best acc: {:.1f}'.format(best_prec))
#     save_checkpoint({
#         'epoch': epoch + 1,
#         'state_dict': model.state_dict(),
#         'best_prec': best_prec,
#         'optimizer': optimizer.state_dict(),
#     }, is_best, fdir)

```

```

[5]: fdir = 'result/'+str(model_name)+'/model_best.pth.tar'

checkpoint = torch.load(fdir)
model.load_state_dict(checkpoint['state_dict'])
criterion = nn.CrossEntropyLoss().cuda()

model.eval()
model.cuda()

prec = validate(testloader, model, criterion)

```

Test: [0/79] Time 0.547 (0.547) Loss 0.5547 (0.5547) Prec 90.625%
(90.625%) * Prec 90.280%

```

[6]: # Hooking
class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in)
    def clear(self):
        self.outputs = []

```

```

[7]: # Save inputs from selected layer
save_output = SaveOutput()
for layer in model.modules():
    if isinstance(layer, nn.Conv2d): # hook input laters of all conv layers
        layer.register_forward_pre_hook(save_output)

dataiter = iter(trainloader)

```

```

images, labels = next(dataiter)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images = images.to(device)
out = model(images)

conv1 = model.layer1[0].conv1
conv2 = model.layer1[0].conv2
bn1 = model.layer1[0].bn1
relu = model.layer1[0].relu
bn2 = model.layer1[0].bn2

# manual calculation with input of the conv1 layer in BasicBlock0 to compute  

the input of the conv1 layer in BasicBlock1

my_input = save_output.outputs[1][0]
my_output = conv1(my_input)
my_output = bn1(my_output)
my_output = relu(my_output)
my_output = conv2(my_output)
my_output = bn2(my_output)
my_output = relu(my_input + my_output)

```

[8]: `print((my_output - save_output.outputs[3][0]).sum())`

```
tensor(0., device='cuda:0', grad_fn=<SumBackward0>)
```