

VGGNet16

October 27, 2025

[1] : # HW

```
# 1. train resnet20 and vgg16 to achieve >90% accuracy
# 2. save your trained model in the result folder
# 3. Restart your jupyter notebook by "Kernel - Restart & Clear Output"
# 4. Load your saved model for vgg16 and validate to see the accuracy
# 5. such as the last part of "[W2S2_example2]_CNN_for_MNIST.ipynb", prehook ↴
#    the input layers of all the conv layers.
# 6. from the first prehooked input, compute to get the second prehooked input.
#    ↴
# 7. Compare your computed second input vs. the prehooked second input.
```

[2] :

```
import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

from models import *    # bring everything in the folder models

global best_prec
use_gpu = torch.cuda.is_available()
print('=> Building model...')

batch_size = 128

model_name = "VGGNet16"
model=VGG16()
```

```

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243, 0.262])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
    ]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
                                          shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
                                         shuffle=False, num_workers=2)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch
# includes "batch_size" data points
# CIFAR10 has 50,000 training data, and 10,000 validation data.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter() ## at the begining of each epoch, this should
# be reset
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time() # measure current time

    for i, (input, target) in enumerate(trainloader):
        # measure data loading time

```

```

data_time.update(time.time() - end) # data loading time

input, target = input.cuda(), target.cuda()

# compute output
output = model(input)
loss = criterion(output, target)

# measure accuracy and record loss
prec = accuracy(output, target)[0]
losses.update(loss.item(), input.size(0))
top1.update(prec.item(), input.size(0))

# compute gradient and do SGD step
optimizer.zero_grad()
loss.backward()
optimizer.step()

# measure elapsed time
batch_time.update(time.time() - end) # time spent to process one batch
end = time.time()

if i % print_freq == 0:
    print('Epoch: [{0}][{1}/{2}]\t'
          'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
          'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
          'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
          'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
              epoch, i, len(trainloader), batch_time=batch_time,
              data_time=data_time, loss=losses, top1=top1))

def validate(val_loader, model, criterion):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

            input, target = input.cuda(), target.cuda()

```

```

# compute output
output = model(input)
loss = criterion(output, target)

# measure accuracy and record loss
prec = accuracy(output, target)[0]
losses.update(loss.item(), input.size(0))
top1.update(prec.item(), input.size(0))

# measure elapsed time
batch_time.update(time.time() - end)
end = time.time()

if i % print_freq == 0: # This line shows how frequently print out
    ↪the status. e.g., i%5 => every 5 batch, prints out
    print('Test: [{0}/{1}]\t'
          'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
          'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
          'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
              i, len(val_loader), batch_time=batch_time, loss=losses,
              top1=top1))

print('* Prec {top1.avg:.3f}% '.format(top1=top1))
return top1.avg

def accuracy(output, target, topk=(1,5)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk) # 5
    batch_size = target.size(0) # 128

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None,
    ↪largest=True, sorted=True) # will output (max value, its index)
    pred = pred.t() # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred)) # "-1": calculate
    ↪automatically

    res = []
    for k in topk: # 1, 5
        correct_k = correct[:k].reshape(-1).float().sum(0) # reshape(-1): make
    ↪a flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size)) # correct: size of
    ↪[maxk, batch_size]
    return res

```

```

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n      ## n is impact factor
        self.count += n
        self.avg = self.sum / self.count

    def save_checkpoint(state, is_best, fdir):
        filepath = os.path.join(fdir, 'checkpoint.pth')
        torch.save(state, filepath)
        if is_best:
            shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

    def adjust_learning_rate(optimizer, epoch):
        """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120 epochs"""
        adjust_list = [150, 225]
        if epoch in adjust_list:
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['lr'] * 0.1

#model = nn.DataParallel(model).cuda()
#all_params = checkpoint['state_dict']
#model.load_state_dict(all_params, strict=False)
#criterion = nn.CrossEntropyLoss().cuda()
#validate(testloader, model, criterion)

```

=> Building model...
 Files already downloaded and verified
 Files already downloaded and verified

[3]: # import matplotlib.pyplot as plt
import numpy as np

```

# # functions to show an image

# def imshow(img):
#     img = img / 2 + 0.5      # unnormalize
#     npimg = img.numpy()
#     plt.imshow(np.transpose(npimg, (1, 2, 0)))
#     plt.show()

# # get some random training images
# dataiter = iter(testloader)
# images, labels = next(dataiter) ## If you run this line, the next data batch
#                                ↴is called subsequently.

# # show images
# imshow(torchvision.utils.make_grid(images))

```

[4]:

```

# # This cell is from the website
# # Comment this section out after training
# lr = 4e-3
# weight_decay = 1e-5
# epochs = 100
# best_prec = 0

# model = model.cuda()
# criterion = nn.CrossEntropyLoss().cuda()
# optimizer = torch.optim.SGD(model.parameters(), lr=lr, momentum=0.9, ↴
#                             weight_decay=weight_decay)
# # weight decay: for regularization to prevent overfitting

# if not os.path.exists('result'):
#     os.makedirs('result')

# fdir = 'result/'+str(model_name)

# if not os.path.exists(fdir):
#     os.makedirs(fdir)

# training_start = time.time()
# for epoch in range(0, epochs):
#     adjust_learning_rate(optimizer, epoch)

#     train(trainloader, model, criterion, optimizer, epoch)

#     # evaluate on test set
#     print("Validation starts")
#     prec = validate(testloader, model, criterion)

```

```

#     # remember best precision and save checkpoint
#     is_best = prec > best_prec
#     best_prec = max(prec,best_prec)
#     print('best acc: {:.1f}'.format(best_prec))
#     save_checkpoint({
#         'epoch': epoch + 1,
#         'state_dict': model.state_dict(),
#         'best_prec': best_prec,
#         'optimizer': optimizer.state_dict(),
#     }, is_best, fdir)
# training_end = time.time() - training_start
# print(training_end)

```

[5]: fdir = 'result/' + str(model_name) + '/model_best.pth.tar'

```

checkpoint = torch.load(fdir)
model.load_state_dict(checkpoint['state_dict'])

```

```
criterion = nn.CrossEntropyLoss().cuda()
```

```
model.eval()
model.cuda()
```

```
prec = validate(testloader, model, criterion)
```

```
Test: [0/79]      Time 0.779 (0.779)      Loss 0.2373 (0.2373)      Prec 93.750%
(93.750%)
* Prec 91.530%
```

[7]: # Hooking

```

class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in)
    def clear(self):
        self.outputs = []

```

[9]: # Save inputs from selected layer

```

save_output = SaveOutput()
for layer in model.modules():
    if isinstance(layer, nn.Conv2d):
        layer.register_forward_pre_hook(save_output)

```

```

dataiter = iter(trainloader)
images, labels = next(dataiter)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images = images.to(device)
out = model(images)

# manual calculation with input of second conv2d layer to output of 1st
↳ maxpool2d layer
conv = model.features[3]
batchnorm = model.features[4]
relu = model.features[5]
maxpool = model.features[6]

my_input = save_output.outputs[1][0]
my_output = conv(my_input)
my_output = batchnorm(my_output)
my_output = relu(my_output)
my_output = maxpool(my_output)
print((my_output - save_output.outputs[2][0]).sum())

```

```
tensor(0., device='cuda:0', grad_fn=<SumBackward0>)
```

[11]:

```

# Save inputs from selected layer
save_output2 = SaveOutput()
for layer in model.modules():
    if isinstance(layer, nn.ReLU) or isinstance(layer, nn.MaxPool2d):
        layer.register_forward_pre_hook(save_output2)

dataiter = iter(trainloader)
images, labels = next(dataiter)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
images = images.to(device)
out = model(images)

# manual calculation with input of first ReLU to input of 1st maxpool2d layer
relu1 = model.features[2]
conv = model.features[3]
batchnorm = model.features[4]
relu2 = model.features[5]

my_input = save_output2.outputs[0][0]
my_output = relu1(my_input)
my_output = conv(my_output)
my_output = batchnorm(my_output)
my_output = relu2(my_output)
print((my_output - save_output2.outputs[2][0]).sum())

```

```
tensor(0., device='cuda:0', grad_fn=<SumBackward0>)
```

Note: I'm confused on the instructions for prehook the input layers of all conv layers and then use grabbed input of the first relu layer. Therefore, I implemented two version, the first one hook to conv layers and use input second conv layer to compute for input for third conv layer (output of 1st maxpool layer). The second one hooked on input layers of relu and maxpool layers, and use input of first ReLU to compute input of 1st maxpool2d layer.